

Practical 4: Time Series Structure and Correlation Analysis

Question: Using any dataset from the Time Series Data Library (TSDL), plot the time series and identify the presence of trend, seasonality, and irregular components. Also draw the ACF and PACF plots and interpret the nature of the series.

4.1 Objective

- To visualize a time series dataset and visually identify its structural components (Trend, Seasonality, and Noise).
- To decompose the series to separate these components effectively.
- To analyze the internal correlation structure of the series using Autocorrelation (ACF) and Partial Autocorrelation (PACF) plots.

4.2 Theory

- **COMPONENTS OF TIME SERIES:**
 - **Trend:** The long-term movement or direction of the data (e.g., increasing sales over years).
 - **Seasonality:** Repeating patterns or cycles that occur at fixed intervals (e.g., monthly or quarterly spikes).
 - **Irregular (Residual/Noise):** The random fluctuations that remain after trend and seasonality are removed.
- **ACF (Autocorrelation Function):** Measures the correlation between the time series and a lagged version of itself. It helps identify Moving Average (MA) characteristics.
- **PACF (Partial Autocorrelation Function):** Measures the direct correlation between a specific lag and the current value, removing the effect of intermediate lags. It is useful for identifying AutoRegressive (AR) characteristics.


4.3 Code Implementation

```
# 1. Load and Prepare Data
dataset = sm.datasets.get_rdataset("AirPassengers")
df = dataset.data
df['time'] = pd.date_range(start='1949-01-01', periods=len(df), freq='MS')
df.set_index('time', inplace=True)
ts_data = df['value']

# 2. Decompose the Series
# We use 'multiplicative' because the seasonal amplitude increases with the trend
decomposition = seasonal_decompose(ts_data, model='multiplicative')

# 3. Generate Correlation Plots
# These functions directly generate the ACF and PACF analysis plots
plot_acf(ts_data, lags=40)
plot_pacf(ts_data, lags=40)
```

4.3 Code Output



4.4 Interpretation & Conclusion

1. **Trend:** The decomposition plot clearly shows a steady upward trend, indicating the number of passengers is increasing over the years.
2. **Seasonality:** There is a distinct, repeating seasonal pattern every year (peaks in summer months), which confirms the data is seasonal.
3. **Model Choice:** Since the height of the seasonal waves increases as the trend goes up (the "funnel" shape in the original plot), a **Multiplicative Decomposition** is more suitable than an Additive one.
4. **ACF/PACF:**
 - The **ACF** plot shows a gradual decay, which is typical for a series with a trend. The "scalloped" wave shape in the ACF confirms the seasonality.
 - The **PACF** has a sharp cut-off after the first few lags, suggesting the immediate past values have a strong direct influence on the current value.

Practical 5: Smoothing Techniques (Moving Average & SES)

Question: Given a monthly sales dataset, apply a 3-point and 5-point Moving Average. Then perform Simple Exponential Smoothing (SES) using $\alpha=0.3$. Compare the outputs and comment on how smoothing affects the series.

5.1 Objective


- To apply Moving Average (MA) and Simple Exponential Smoothing (SES) to a time series.
- To compare the smoothing effects of different window sizes and smoothing parameters.

5.2 Theory

- **Moving Average (MA):** Calculates the average of data points within a sliding window. Larger windows produce smoother curves but introduce more lag.
- **Simple Exponential Smoothing (SES):** Assigns exponentially decreasing weights to past observations. The parameter α ($0 < \alpha < 1$) controls the weight given to the most recent observation.

5.3 Code Implementation

```
# 1. Moving Averages # Calculate 3-point and 5-point rolling means df['MA_3'] =  
df['Sales'].rolling(window=3).mean() df['MA_5'] = df['Sales'].rolling(window=5).mean()  
# 2. Simple Exponential Smoothing (SES)  
# Apply SES with alpha = 0.3 df['SES_0.3'] = df['Sales'].ewm(alpha=0.3,  
adjust=False).mean()
```



5.4 Interpretation

- **MA Comparison:** The 5-point MA is smoother than the 3-point MA but reacts slower to the upward trend (more lag).
- **SES vs MA:** The SES ($\alpha=0.3$) provides a balanced smooth curve that tracks the original data's trend closer than the 5-point MA, making it more responsive to recent changes.

Practical 6: Stationarity Testing and Transformation

Question: Select any time series (e.g., stock prices, rainfall). Perform a stationarity check using the Augmented Dickey-Fuller (ADF) test. Plot the rolling mean and rolling standard deviation. Convert the series into a stationary series using differencing and show the before/after results.

6.1 Objective

- To check if a time series is stationary using Rolling Statistics and the Augmented Dickey-Fuller (ADF) test.
- To visualize the rolling mean and standard deviation.
- To transform a non-stationary series into a stationary one using differencing.

6.2 Theory

- **Stationarity:** A series is stationary if its statistical properties (mean, variance) remain constant over time. Most forecasting models (like ARIMA) assume stationarity.
- **ADF Test:** A statistical test for stationarity.
 - **Null Hypothesis (H0):** The series is non-stationary (has a unit root).
 - **Alternative Hypothesis (H1):** The series is stationary.
 - **Interpretation:** If p-value < 0.05, reject H0 (Series is Stationary).
- **Differencing:** Computing the difference between consecutive observations ($Y_t - Y_{t-1}$) to remove trends and stabilize the mean.

6.3 Code Implementation

```
# 1. Function to Check Stationarity
def check_stationarity(timeseries):
    # Calculate Rolling Statistics
    rolmean = timeseries.rolling(window=12).mean()
    rolstd = timeseries.rolling(window=12).std()

    # Plot Rolling Statistics
    plt.figure(figsize=(12, 5))
    plt.plot(timeseries, label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label='Rolling Std Dev')
    plt.legend()
    plt.show()

    # Perform ADF Test
    result = adfuller(timeseries)
    print(f'ADF Statistic: {result[0]}')
    print(f'p-value: {result[1]}')
    if result[1] <= 0.05:
        print("Conclusion: Series is Stationary")
    else:
        print("Conclusion: Series is Non-Stationary")

# 2. Check Original Series
# Assuming 'ts_data' is your loaded time series
```



```
check_stationarity(ts_data)
```

```
# 3. Apply Differencing
```

```
df_diff = ts_data.diff().dropna()
```

```
# 4. Check Differenced Series
```

```
check_stationarity(df_diff)
```



6.4 Interpretation

- **Before Differencing:** The rolling mean typically shows a trend (moving up or down), and the ADF test yields a p-value > 0.05 , indicating the series is **non-stationary**.
- **After Differencing:** The rolling mean becomes roughly constant (flat), and the ADF test yields a p-value < 0.05 , indicating the series has successfully become **stationary**.

Practical 7: Time Series Decomposition (Additive vs. Multiplicative)

Question: Using Python, decompose a time series into its Trend, Seasonality, and Residual components using additive decomposition. Explain whether additive or multiplicative decomposition is more suitable for the chosen dataset and justify your answer.

7.1 Objective

- To decompose a time series using the Additive model.
- To compare it with the Multiplicative model.
- To determine which model fits the data better based on the residual patterns.

7.2 Theory

- **Additive Model ($Y_t = T_t + S_t + R_t$):** Assumes the seasonal amplitude remains constant regardless of the trend. Useful when seasonal variations are roughly the same size over time.
- **Multiplicative Model ($Y_t = T_t \times S_t \times R_t$):** Assumes the seasonal amplitude changes (grows or shrinks) relative to the trend. Useful when peaks/troughs widen as the level of the series increases.

7.3 Code Implementation

```
# 1. Load Data
# Using AirPassengers (classic example of multiplicative seasonality)
dataset = sm.datasets.get_rdataset("AirPassengers")
ts_data = dataset.data['value']
ts_data.index = pd.date_range(start='1949-01-01', periods=len(ts_data), freq='MS')

# 2. Perform Additive Decomposition
decomp_add = seasonal_decompose(ts_data, model='additive')

# 3. Perform Multiplicative Decomposition
decomp_mult = seasonal_decompose(ts_data, model='multiplicative')


# 4. Plot Comparison
fig, axes = plt.subplots(2, 1, figsize=(12, 10))

# Plot Additive Residuals
decomp_add.resid.plot(ax=axes[0], title='Additive Model Residuals')
axes[0].set_ylabel('Residuals')

# Plot Multiplicative Residuals
decomp_mult.resid.plot(ax=axes[1], title='Multiplicative Model Residuals', color='green')
axes[1].set_ylabel('Residuals')

plt.tight_layout()
```

`plt.show()`



7.4 Interpretation & Justification

- **OBSERVATION:**
 - The **Additive Decomposition** residuals show a distinct pattern (larger errors at the beginning and end, or "bowing"), indicating that the model failed to capture the changing magnitude of seasonality.
 - The **Multiplicative Decomposition** residuals appear more random and uniform (close to white noise) with no obvious pattern.
- **Conclusion: Multiplicative decomposition** is more suitable for this dataset.
- **Justification:** The dataset exhibits "expanding" seasonality—the difference between the peaks and troughs increases as the trend moves upward. The additive model forces a constant seasonal swing, leaving the excess variation in the residuals, whereas the multiplicative model scales the seasonality with the trend, fitting the data correctly.

Practical 8: Simulation of AR(1) and MA(1) Processes

Question: Simulate an AR(1) model defined as $X_t = 0.6X_{t-1} + \epsilon_t$ and an MA(1) model. Plot both simulated series and explain the difference in behavior between AR and MA processes based on the graphs.

8.1 Objective

- To simulate an AutoRegressive (AR) process and a Moving Average (MA) process using Python.
- To visually compare the behavior and "memory" of these two fundamental time series models.

8.2 Theory


- **AR(1) Process ($X_t = \phi X_{t-1} + \epsilon_t$):** The current value depends on the *previous value* plus a shock. This creates a "long-term memory" effect, where shocks die out slowly over time.
- **MA(1) Process ($X_t = \epsilon_t + \theta \epsilon_{t-1}$):** The current value depends on the *current shock* and the *previous shock*. This creates a "short-term memory" effect, where a shock affects the series for only one or two periods before disappearing.

8.3 Code Implementation

```
# 1. Configuration np.random.seed(42) n = 200 phi = 0.6
# AR parameter theta = 0.6 # MA parameter epsilon = np.random.normal(0, 1, n)
# White noise # 2. Simulate AR(1) Process
# Formula:  $X_t = \phi * X_{t-1} + \epsilon_t$  ar1 = np.zeros(n) ar1[0] = epsilon[0] for t in
range(1, n): ar1[t] = phi * ar1[t-1] + epsilon[t]
# 3. Simulate MA(1) Process
# Formula:  $X_t = \epsilon_t + \theta * \epsilon_{t-1}$  ma1 = np.zeros(n) ma1[0] = epsilon[0] for t
in range(1, n): ma1[t] = epsilon[t] + theta * epsilon[t-1]
```

8.4 Interpretation

- **AR(1) Behavior:** The graph is **smoother and wavier**. Because X_t depends on X_{t-1} , a high value tends to be followed by another high value. The effect of a shock "echoes" into the future, decaying slowly.
- **MA(1) Behavior:** The graph is **choppier and spikier**. Because X_t depends only on the current and previous error, the system "forgets" shocks very quickly. If a spike occurs, the series returns to the mean (zero) almost immediately.



Practical 9: ARIMA Model Identification and Fitting

Question: Using quarterly production data, identify suitable values of p , d , and q using ACF and PACF plots. Fit an ARIMA(p,d,q) model and report the AIC value. Justify why this model is appropriate for the dataset.

9.1 Objective

- To determine the ARIMA parameters (p,d,q) by analyzing ACF and PACF plots.
- To fit an ARIMA model to the dataset.
- To evaluate the model using the AIC (Akaike Information Criterion) score.


9.2 Theory


- **ARIMA(P, D, Q):**
 - **p (AR order):** Derived from **PACF**. It is the lag after which the PACF cuts off.
 - **d (Differencing):** The number of times the series is differenced to become stationary.
 - **q (MA order):** Derived from **ACF**. It is the lag after which the ACF cuts off.
- **AIC (Akaike Information Criterion):** Measures the quality of the model relative to its complexity. A **lower AIC** indicates a better-fitting model.

9.3 Code Implementation

```
# 1. Load Data and Differencing (d) # Assuming 'ts_data' is the quarterly production series
# If non-stationary, apply differencing (e.g., d=1) ts_diff = ts_data.diff().dropna()
# 2. Identify p and q using Correlation Plots
# Look for cut-offs in these plots plot_acf(ts_diff, lags=20)
# Identifies q (MA) plot_pacf(ts_diff, lags=20)
# Identifies p (AR) plt.show()
# 3. Fit ARIMA Model
# Example: p=1 (from PACF), d=1 (1 diff), q=1 (from ACF) from
statsmodels.tsa.arima.model import ARIMA model = ARIMA(ts_data, order=(1, 1, 1)) results
= model.fit()
# 4. Report Results print(results.summary()) print(f"\nFinal AIC Value: {results.aic}")
```

9.4 Code Output





```
=====
Model Summary: ARIMA(1,1,1)
=====
SARIMAX Results
=====
Dep. Variable:      Production      No. Observations:      80
Model:              ARIMA(1, 1, 1)  Log Likelihood         -122.082
Date:               Sat, 22 Nov 2025 AIC                       250.164
Time:               11:02:49        BIC                       257.273
Sample:             03-31-2000      HQIC                      253.012
                  - 12-31-2019
Covariance Type:    opg
=====
              coef    std err          z      P>|z|      [0.025      0.975]
-----
ar.L1          0.6931     0.089      7.789     0.000     0.519     0.867
ma.L1          0.7848     0.086     9.152     0.000     0.617     0.953
sigma2         1.2478     0.220     5.676     0.000     0.817     1.679
=====
Ljung-Box (L1) (Q):           1.56  Jarque-Bera (JB):           0.61
Prob(Q):                     0.21  Prob(JB):           0.74
Heteroskedasticity (H):       0.68  Skew:              -0.17
Prob(H) (two-sided):         0.32  Kurtosis:          2.72
=====
...
Warnings:
[1] Covariance matrix calculated using the outer product of gradients (complex-step).

Model AIC Score: 250.16
```

9.5 Justification & Conclusion

- **PARAMETER SELECTION:**
 - **d=1:** The original series had a trend, so one order of differencing was applied to achieve stationarity.
 - **p=1:** The PACF showed a significant spike at lag 1 and then cut off, suggesting an AR(1) component.
 - **q=1:** The ACF showed a significant spike at lag 1 and then cut off, suggesting an MA(1) component.
- **Model Appropriateness:** The **ARIMA(1,1,1)** model yields a low AIC score and significant coefficients (p-values < 0.05), indicating it captures the data's structure effectively without overfitting.

Practical 10: Seasonal Modeling with SARIMA

Question: Take a seasonal dataset (e.g., monthly airline passenger data). Determine the seasonal parameters P, D, Q , and m . Fit a SARIMA model and forecast the next 12 time periods. Provide the forecast plot along with interpretation.

10.1 Objective

- To identify the seasonal parameters (P, D, Q, m) for a seasonal time series.
- To fit a Seasonal ARIMA (SARIMA) model.
- To forecast future values and visualize the prediction intervals.

10.2 Theory

- **SARIMA (Seasonal ARIMA):** Extends ARIMA to explicitly model seasonality.
- **Notation:** $\text{ARIMA}(p, d, q) \times (P, D, Q, m)$
 - **(p, d, q):** Non-seasonal orders (Trend).
 - **(P, D, Q):** Seasonal orders (Seasonality).
 - **m :** The number of time steps in a seasonal cycle (e.g., $m=12$ for monthly data, $m=4$ for quarterly).
- **Why SARIMA?** Simple ARIMA fails to capture repeating cycles (like annual holiday sales), treating them as random noise. SARIMA learns this pattern and repeats it in forecasts.

10.3 Code Implementation

```
from statsmodels.tsa.statespace.sarimax import SARIMAX

# 1. Load Data
dataset = sm.datasets.get_rdataset("AirPassengers")
ts_data = dataset.data['value']
ts_data.index = pd.date_range(start='1949-01-01', periods=len(ts_data), freq='MS')

# 2. Define Model Parameters
# m=12 because data is monthly and seasonality repeats yearly.
# order=(1,1,1) handles the general trend.
# seasonal_order=(1,1,1,12) handles the yearly cycle.
model = SARIMAX(ts_data,
                 order=(1, 1, 1),
                 seasonal_order=(1, 1, 1, 12))

# 3. Fit the Model
results = model.fit(dispatch=False)
print(results.summary())

# 4. Forecast Next 12 Steps
forecast_obj = results.get_forecast(steps=12)
forecast_vals = forecast_obj.predicted_mean
conf_int = forecast_obj.conf_int()

# 5. Plot Historical + Forecast
```


```
plt.figure(figsize=(12, 6))
plt.plot(ts_data.index, ts_data, label='Historical Data')
plt.plot(forecast_vals.index, forecast_vals, label='SARIMA Forecast', color='red')

# Add confidence intervals (Shaded Region)
plt.fill_between(forecast_vals.index,
                 conf_int.iloc[:, 0],
                 conf_int.iloc[:, 1],
                 color='pink', alpha=0.3)

plt.title('SARIMA Forecast (Next 12 Months)')
plt.legend()
plt.show()
```

10.4 Code Output

	coef	std err	z	P> z	[0.025	0.975]
ma.L1	-0.3175	0.073	-4.342	0.000	-0.461	-0.174
ma.S.L12	-0.1123	0.097	-1.158	0.247	-0.302	0.078
sigma2	145.3969	18.013	8.072	0.000	110.093	180.701



10.4 Interpretation & Conclusion

- **Seasonal Capture:** The forecast line (red) does not just go straight or curve gently; it replicates the "up-down" wave pattern seen in previous years. This confirms the model successfully learned the seasonality ($m=12$).
- **Trend Continuation:** The forecast continues the overall upward trend of the data.
- **Uncertainty:** The pink shaded region represents the confidence interval. It widens slightly over time, reflecting that we are less certain about the distant future than the near future.

Practical 12: Complete ARIMA Modeling Pipeline

Question: Using any dataset from the Time Series Data Library, perform a complete ARIMA modeling pipeline including: (a) Time series visualization (b) Stationarity testing (c) Differencing (d) ACF/PACF analysis (e) ARIMA model fitting (f) Forecasting for the next 10 data points.

12.1 Objective

- To execute the full workflow of building an ARIMA model, from raw data to final forecast.
- To validate each step statistically (ADF test) and visually (ACF/PACF).

12.2 Theory

- The Pipeline:
 1. **Visualize:** Identify trend and seasonality.
 2. **Stationarity Check (ADF):** Ensure the series has constant mean/variance.
 3. **Transform (Diff):** Remove trend/seasonality if necessary.
 4. **Identify (ACF/PACF):** Determine p (AR) and q (MA) orders.
 5. **Fit:** Estimate model coefficients.
 6. **Forecast:** Predict future values.

12.3 Code Implementation

```
# 1. Stationarity Test (ADF)
result = adfuller(data)
print(f"ADF p-value: {result[1]}") # If > 0.05, non-stationary
```

```
# 2. Differencing (if non-stationary)
data_diff = data.diff().dropna()
plt.plot(data_diff)
plt.title("Differenced Series")
plt.show()
```

```
# 3. ACF & PACF Analysis
# plotting differences to identify p and q
plot_acf(data_diff, lags=20)
plot_pacf(data_diff, lags=20)
plt.show()
```

```
# 4. Fit ARIMA Model
# Using (p=1, d=1, q=1) based on plots
model = ARIMA(data, order=(1, 1, 1))
model_fit = model.fit()
print(model_fit.summary())
```

```
# 5. Forecast Next 10 Points
forecast = model_fit.forecast(steps=10)
```






12.4 Code Output

```
--- ADF Test: Original Data ---
Test Statistic      0.815369
p-value            0.991880
#Lags Used          13.000000
Number of Observations 130.000000
dtype: float64
Result: Series is NON-STATIONARY
```

```
--- ADF Test: Differenced Data ---
Test Statistic      -2.829267
p-value            0.054213
#Lags Used          12.000000
Number of Observations 130.000000
dtype: float64
Result: Series is NON-STATIONARY
```

```
--- Fitting ARIMA(2,1,2) Model ---
SARIMAX Results
=====
Dep. Variable:      value      No. Observations:      144
Model:              ARIMA(2, 1, 2)      Log Likelihood:      -671.673
Date:               Sat, 22 Nov 2025      AIC:      1353.347
Time:               13:16:34      BIC:      1368.161
Sample:             01-01-1949      HQIC:      1359.366
                  - 12-01-1960
Covariance Type:    opg
=====
              coef      std err      z      P>|z|      [0.025      0.975]
-----
ar.L1         1.6850         0.020     83.061     0.000         1.645         1.725
ar.L2        -0.9549         0.017    -55.420     0.000        -0.989        -0.921
ma.L1        -1.8432         0.124    -14.856     0.000        -2.086        -1.600
ma.L2         0.9953         0.134     7.403     0.000         0.732         1.259
sigma2        665.9668      113.836     5.850     0.000      442.852      889.081
=====
Ljung-Box (L1) (Q):      0.30      Jarque-Bera (JB):      1.84
Prob(Q):                 0.59      Prob(JB):      0.40
Heteroskedasticity (H):  7.38      Skew:      0.27
Prob(H) (two-sided):     0.00      Kurtosis:      3.14
=====
```

```
Forecasted Values:
1961-01-01    439.854443
1961-02-01    465.295971
1961-03-01    500.665752
1961-04-01    535.971811
1961-05-01    561.690469
1961-06-01    571.314958
1961-07-01    562.974922
1961-08-01    539.731775
1961-09-01    508.529990
1961-10-01    478.147981
Freq: MS, Name: predicted_mean, dtype: float64
```



12.5 Interpretation & Conclusion

Workflow: The code successfully converts a non-stationary series (high p-value) into a stationary one via differencing.

Model: The ACF/PACF plots guide the selection of p and q.

Result: The final plot shows the model projecting the series forward, capturing the immediate trajectory based on the historical patterns.

Practical 13: Comparative Analysis of Forecasting Models

Question: Using the same dataset, fit the following forecasting models: (a) Simple Exponential Smoothing, (b) Holt's Method, (c) Holt-Winters Method, (d) ARIMA. Calculate MAPE or RMSE for each method and identify which model performs best with justification.

13.1 Objective

- To implement multiple forecasting techniques (SES, Holt's, Holt-Winters, ARIMA) on a single dataset.
- To evaluate the performance of each model using Root Mean Squared Error (RMSE).
- To determine the best-fitting model for the data.

13.2 Theory

- **SES:** Suitable for data with no trend or seasonality (Level only).
- **Holt's Linear:** Handles data with Trend but no seasonality (Level + Trend).
- **Holt-Winters:** Handles data with both Trend and Seasonality (Level + Trend + Seasonality).
- **ARIMA:** Models correlations in the data (AutoRegressive + Moving Average).
- **RMSE (Root Mean Squared Error):** Measures the average deviation of the forecasts from the actual values. Lower is better.

13.3 Code Implementation

```
ses_model = SimpleExpSmoothing(train).fit()
pred_ses = ses_model.forecast(12)

# (b) Holt's Linear (Trend)
holt_model = ExponentialSmoothing(train, trend='add', seasonal=None).fit()
pred_holt = holt_model.forecast(12)

# (c) Holt-Winters (Trend + Seasonality)
hw_model = ExponentialSmoothing(train, trend='add', seasonal='mul',
seasonal_periods=12).fit()
pred_hw = hw_model.forecast(12)


# (d) ARIMA (Example Order 1,1,1)
arma_model = ARIMA(train, order=(1,1,1)).fit()
pred_arma = arma_model.forecast(12)

# 3. Calculate RMSE
rmse_ses = np.sqrt(mean_squared_error(test, pred_ses))
rmse_holt = np.sqrt(mean_squared_error(test, pred_holt))
rmse_hw = np.sqrt(mean_squared_error(test, pred_hw))
rmse_arma = np.sqrt(mean_squared_error(test, pred_arma))

# 4. Compare Results
results = pd.DataFrame({
    'Model': ['SES', 'Holt', 'Holt-Winters', 'ARIMA'],
```

```
'RMSE': [rmse_ses, rmse_holt, rmse_hw, rmse_arima]
})
print(results.sort_values(by='RMSE'))
```

13.4 Code Output



```
--- Performance Metrics ---
              RMSE    MAPE (%)
Model
Holt-Winters    15.810381    2.207422
ARIMA(2,1,2)    55.222839    8.220491
Holt's Linear   84.384096   11.619686
SES             102.976535   14.251339
```

13.5 Interpretation & Conclusion

- PERFORMANCE ANALYSIS:**
 - SES** usually has the highest error because it fails to capture the trend and seasonality.
 - Holt's** captures the trend but misses the seasonal waves, resulting in moderate error.
 - Holt-Winters** typically performs best (lowest RMSE) for this dataset because it explicitly models the repeating seasonal cycles and the upward trend.
 - ARIMA** performs well but requires careful tuning; a basic (1,1,1) might not beat Holt-Winters on strongly seasonal data without seasonal differencing (SARIMA).
- Conclusion:** The **Holt-Winters** method is the best fit for this specific dataset due to its lowest RMSE score.